



Java Performance Analysis on Linux with Flame Graphs

Brendan Gregg Senior Performance Architect





Complete Visibility




NETFLIX Cloud

- Tens of thousands of AWS EC2 instances
- Mostly Java (Oracle JVM)




The Problem with Profilers

Java Profilers




Java Profilers

- **Visibility**
 - Java method execution
 - Object usage
 - GC logs
 - Custom Java context
- **Typical problems:**
 - Sampling often happens at safety/yield points (skew)
 - Method tracing has massive observer effect
 - Misidentifies RUNNING as on-CPU (e.g., epoll)
 - Doesn't include or profile GC or JVM CPU time
 - Tree views not quick (proportional) to comprehend
- **Inaccurate (skewed) and incomplete profiles**




System Profilers




System Profilers

- **Visibility**
 - JVM (C++)
 - GC (C++)
 - libraries (C)
 - kernel (C)
- Typical problems (x86):
 - Stacks missing for Java
 - Symbols missing for Java methods
- Other architectures (e.g., SPARC) have fared better
- Profile everything except Java




Workaround

- Capture both:



Java




System

- An improvement, but system stacks are missing Java context, and therefore hard to interpret


Solution

Java Mixed-Mode Flame Graph



Solution

- Fix system profiling, see everything:
 - Java methods
 - JVM (C++)
 - GC (C++)
 - libraries (C)
 - kernel (C)
 - Other apps
- Minor Problems:
 - 0-3% CPU overhead to enable frame pointers (usually <1%).
 - Symbol dumps can consume a burst of CPU
- **Complete and accurate (asynchronous) profiling**




Saving 13M CPU Minutes Per Day



<http://techblog.netflix.com/2016/04/saving-13-million-computational-minutes.html>

System Example

Exception handling consuming CPU



Profiling GC


GC internals, visualized:



CPU Profiling

CPU Profiling

- Record stacks at a timed interval: simple and effective
 - Pros: Low (deterministic) overhead
 - Cons: Coarse accuracy, but usually sufficient



Stack Traces

- A code path snapshot. e.g., from jstack(1):

```
$ jstack 1819
[...]
"main" prio=10 tid=0x00007ff304009000
nid=0x7361 runnable [0x00007ff30d4f9000]
    java.lang.Thread.State: RUNNABLE
        at Func_abc.func_c(Func_abc.java:6)
        at Func_abc.func_b(Func_abc.java:16)
        at Func_abc.func_a(Func_abc.java:23)
        at Func_abc.main(Func_abc.java:27)
```

running
parent
g.parent
g.g.parent

System Profilers

- Linux
 - perf_events (aka "perf")
- Oracle Solaris
 - DTrace
- OS X
 - Instruments
- Windows
 - XPerf, WPA (which now has flame graphs!)
- And many others...

Linux perf_events


- Standard Linux profiler
 - Provides the `perf` command (multi-tool)
 - Usually pkg added by `linux-tools-common`, etc.
- Many event sources:
 - Timer-based sampling
 - Hardware events
 - Tracepoints
 - Dynamic tracing
- Can sample stacks of (almost) everything on CPU
 - Can miss hard interrupt ISRs, but these should be near-zero. They can be measured if needed (I wrote my own tools)

perf Profiling


```
# perf record -F 99 -ag -- sleep 30
[ perf record: Woken up 9 times to write data ]
[ perf record: Captured and wrote 2.745 MB perf.data (-119930 samples) ]
# perf report -n -stdio
[...]
# Overhead          Samples  Command           Shared Object          Symbol
# .....  .....
#
20.42%          605      bash   [kernel.kallsyms]  [k] xen_hypcall_xen_version
|
--- xen_hypcall_xen_version
    check_events
    |
    ---44.13%--- syscall_trace_enter
        tracesys
        |
        ---35.58%--- __GI__libc_fcntl
                    |
                    ---65.26%--- do_redirection_internal
                                do_redirections
                                execute_builtin_or_function
                                execute_simple_command
[... ~13,000 lines truncated ...]
```

call tree summary

Full perf report Output




... as a Flame Graph




Flame Graphs

- Flame Graphs:
 - **x-axis**: alphabetical stack sort, to maximize merging
 - **y-axis**: stack depth
 - **color**: random (default), or a dimension
 - Currently made from Perl + SVG + JavaScript
 - Multiple d3 versions are being developed
 - References:
 - <http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html>
 - <http://queue.acm.org/detail.cfm?id=2927301>
 - "The Flame Graph" CACM, June 2016
 - Easy to make
 - Converters for many profilers




Flame Graph Interpretation



Flame Graph Interpretation (1/3)


Top edge shows who is running on-CPU,
and how much (width)



Flame Graph Interpretation (2/3)

Top-down shows ancestry


e.g., from g():



Flame Graph Interpretation (3/3)


Widths are proportional to presence in samples

e.g., comparing `b()` to `h()` (incl. children)




Mixed-Mode Flame Graphs

- Hues:
 - green == Java
 - aqua == Java (inlined)
 - if included
 - red == system
 - yellow == C++
- Intensity:
 - Randomized to differentiate frames
 - Or hashed on function name




Differential Flame Graphs

- Hues:
 - red == more samples
 - blue == less samples
- Intensity:
 - Degree of difference
- Compares two profiles
- Can show other metrics: e.g., CPI
- Other types exist
 - flamegraphdiff



Flame Graph Search

- Color: magenta to show matched frames



Flame Charts

- Final note: these are useful, but are not flame *graphs*



- Flame **charts**: x-axis is time
- Flame **graphs**: x-axis is population (maximize merging)

Stack Tracing


Broken Java Stacks on x86

- These stacks are 1 or 2 levels deep, with junk values
- On x86 (x86_64), hotspot uses the frame pointer register (RBP) as general purpose
- This "compiler optimization" breaks (RBP-based) stack walking
- *Once upon a time*, x86 had fewer registers, and this made more sense
- gcc provides **-fno-omit-frame-pointer** to avoid doing this, but the JVM had no such option...

```
# perf record -F 99 -a -g - sleep 30
# perf script
[...]
java 4579 cpu-clock:
    7f417908c10b [unknown] (/tmp/perf-4458.map)

java 4579 cpu-clock:
    7f41792fc65f [unknown] (/tmp/perf-4458.map)
    a2d53351ff7da603 [unknown] ([unknown])
[...]
```

... as a Flame Graph



Fixing Stack Walking

Possibilities:

- A. Fix frame pointer-based stack walking (the default)
 - Pros: simple, supported by many tools
 - Cons: might cost a little extra CPU
- B. Use a custom walker (likely needing kernel support)
 - Pros: full stack walking (incl. inlining) & arguments
 - Cons: custom kernel code, can cost more CPU when in use
- C. Try libunwind and DWARF
 - Even feasible with JIT?

Our current preference is (A)

-XX:+PreserveFramePointer

- I hacked OpenJDK x86_64 to support frame pointers
 - Taking RBP out of register pools, and adding function prologues. It worked, I shared the patch.
 - It became **JDK-8068945** for JDK 9 and **JDK-8072465** for JDK 8
- Zoltán Majó (Oracle) rewrote it, and it is now:
 - -XX:+PreserveFramePointer in **JDK 9** and **JDK 8 u60b19**
 - Thanks to Zoltán, Oracle, and the other hotspot engineers for helping get this done!
- It might cost 0 – 3% CPU, depending on workload


Fixed Java Stacks

```
# perf script
[...]
java 4579 cpu-clock:
    7f417908c10b [unknown] (/tmp/...
java 4579 cpu-clock:
    7f41792fc65f [unknown] (/tmp/...
    a2d53351ff7da603 [unknown] ([unkn...
[...]
```




```
# perf script
[...]
java 8131 cpu-clock:
    7fff76f2dce1 [unknown] ([vdso])
    7fd3173f7a93 os::javaTimeMillis() (/usr/lib/jvm...
    7fd301861e46 [unknown] (/tmp/perf-8131.map)
    7fd30184def8 [unknown] (/tmp/perf-8131.map)
    7fd30174f544 [unknown] (/tmp/perf-8131.map)
    7fd30175d3a8 [unknown] (/tmp/perf-8131.map)
    7fd30166d51c [unknown] (/tmp/perf-8131.map)
    7fd301750f34 [unknown] (/tmp/perf-8131.map)
    7fd3016c2280 [unknown] (/tmp/perf-8131.map)
    7fd301b02ec0 [unknown] (/tmp/perf-8131.map)
    7fd3016f9888 [unknown] (/tmp/perf-8131.map)
    7fd3016ece04 [unknown] (/tmp/perf-8131.map)
    7fd30177783c [unknown] (/tmp/perf-8131.map)
    7fd301600aa8 [unknown] (/tmp/perf-8131.map)
    7fd301a4484c [unknown] (/tmp/perf-8131.map)
    7fd3010072e0 [unknown] (/tmp/perf-8131.map)
    7fd301007325 [unknown] (/tmp/perf-8131.map)
    7fd301007325 [unknown] (/tmp/perf-8131.map)
    7fd3010004e7 [unknown] (/tmp/perf-8131.map)
    7fd3171df76a JavaCalls::call_helper(JavaValue*,...
    7fd3171dce44 JavaCalls::call_virtual(JavaValue*...
    7fd3171dd43a JavaCalls::call_virtual(JavaValue*...
    7fd31721b6ce thread_entry(JavaThread*, Thread*)...
    7fd3175389e0 JavaThread::thread_main_inner() (/...
    7fd317538cb2 JavaThread::run() (/usr/lib/jvm/nf...
    7fd3173f6f52 java_start(Thread*) (/usr/lib/jvm/...
    7fd317a7e182 start_thread (/lib/x86_64-linux-gn...
```

Fixed Stacks Flame Graph



Stack Depth

- perf had a 127 frame limit
- Now tunable in Linux 4.8
 - `sysctl -w kernel.perf_event_max_stack=512`
 - Thanks Arnaldo Carvalho de Melo!



Symbols

Fixing Symbols

- For JIT'd code, Linux perf already looks for an externally provided symbol file: /tmp/perf-PID.map, and warns if it doesn't exist

```
# perf script
Failed to open /tmp/perf-8131.map, continuing without symbols
[...]
java 8131 cpu-clock:
    7fff76f2dce1 [unknown] ([vdso])
    7fd3173f7a93 os:::javaTimeMillis() (/usr/lib/jvm...
    7fd301861e46 [unknown] (/tmp/perf-8131.map)
[...]
```


- This file can be created by a Java agent

Java Symbols for perf


- perf-map-agent
 - <https://github.com/jrudolph/perf-map-agent>
 - Agent attaches and writes the /tmp file on demand (previous versions attached on Java start, wrote continually)
 - Thanks Johannes Rudolph!
- Use of a /tmp symbol file
 - Pros: simple, can be low overhead (snapshot on demand)
 - Cons: stale symbols
- Using a symbol logger with perf instead
 - Stephane Eranian contributed this to perf
 - See lkml for "perf: add support for profiling jitted code"

```
# perf script
java 14025 [017] 8048.157085: cpu-clock: 7fd781253265 Ljava/util/
HashMap;::get (/tmp/perf-12149.map)
[...]
```


fixed
symbols



Stacks & Symbols




Stacks & Symbols (zoom)



Inlining

- Many frames may be missing (inlined)
 - Flame graph may still make enough sense
 - Inlining can be tuned
 - -XX:-Inline to disable, but can be 80% slower!
 - -XX:MaxInlineSize and -XX:InlineSmallCode can be tuned a little to reveal more frames
 - Can even improve performance!
 - perf-map-agent can un-inline (unfoldall)
 - Adds inlined frames to symbol dump
 - flamegraph.pl --color=java will color these aqua –
 - Thanks Johannes Rudolph, T Jake Luciani, and Nitsan Wakart



Instructions

Instructions

1. Check Java version
2. Install perf-map-agent
3. Set -XX:+PreserveFramePointer
4. Profile Java
5. Dump symbols
6. Generate Mixed-Mode Flame Graph

Note these are unsupported: use at your own risk

Reference: <http://techblog.netflix.com/2015/07/java-in-flames.html>

1. Check Java Version

- Need JDK8u60 or better
 - for -XX:+PreserveFramePointer

```
$ java -version
java version "1.8.0_60"
Java(TM) SE Runtime Environment (build 1.8.0_60-b27)
Java HotSpot(TM) 64-Bit Server VM (build 25.60-b23, mixed mode)
```

- Upgrade Java if necessary

2. Install perf-map-agent

- Check <https://github.com/jrudolph/perf-map-agent> for the latest instructions. e.g.:

```
$ sudo bash
# apt-get install -y cmake
# export JAVA_HOME=/usr/lib/jvm/java-8-oracle
# cd /usr/lib/jvm
# git clone --depth=1 https://github.com/jrudolph/perf-map-agent
# cd perf-map-agent
# cmake .
# make
```

3. Set -XX:+PreserveFramePointer

- Needs to be set on Java startup
- Check it is enabled (on Linux):

```
$ ps wwp `pgrep -n java` | grep PreserveFramePointer
```

or

```
$ jcmd `pgrep -n java` VM.flags | grep PreserveFramePointer
```

- Measure overhead (should be small)

4. Profile Java

- Using Linux perf_events to profile all processes, at 99 Hertz, for 30 seconds (as root):

```
# perf record -F 99 -a -g -- sleep 30
```

- Just profile one PID (broken on some older kernels):

```
# perf record -F 99 -p PID -g -- sleep 30
```

- These create a perf.data file

5. Dump Symbols

- See perf-map-agent docs for updated usage
- e.g., as the same user as java:

```
$ cd /usr/lib/jvm/perf-map-agent/out
$ java -cp attach-main.jar:$JAVA_HOME/lib/tools.jar \
    net.virtualvoid.perf.AttachOnce PID
```

- perf-map-agent contains helper scripts. I wrote my own:
 - <https://github.com/brendangregg/Misc/blob/master/java/jmaps>
- Dump symbols quickly after perf record to minimize stale symbols. How I do it:

```
# perf record -F 99 -a -g -- sleep 30; jmaps
```

6. Generate a Mixed-Mode Flame Graph

- Using my FlameGraph software:

```
# perf script > out.stacks01
# git clone --depth=1 https://github.com/brendangregg/FlameGraph
# cat out.stacks01 | ./FlameGraph/stackcollapse-perf.pl | \
    ./FlameGraph/flamegraph.pl --color=java --hash > flame01.svg
```

- perf script reads perf.data with /tmp/*.map
- out.stacks01 is an intermediate file; can be handy to keep
- Finally open flame01.svg in a browser
- Check for newer flame graph implementations (e.g., d3)

Optimizations

- Linux 2.6+, via perf.data and perf script:

```
git clone --depth 1 https://github.com/brendangregg/FlameGraph
cd FlameGraph
perf record -F 99 -a -g -- sleep 30
perf script | ./stackcollapse-perf.pl | ./flamegraph.pl > perf.svg
```


- Linux 4.5+ can use folded output
 - Skips the CPU-costly stackcollapse-perf.pl step; see:
<http://www.brendangregg.com/blog/2016-04-30/linux-perf-folded.html>
- Linux 4.9+, via BPF:

```
git clone --depth 1 https://github.com/brendangregg/FlameGraph
git clone --depth 1 https://github.com/iovisor/bcc
./bcc/tools/profile.py -dF 99 30 | ./FlameGraph/flamegraph.pl > perf.svg
```


- Most efficient: no perf.data file, summarizes in-kernel

Linux Profiling Optimizations


Linux 2.6



Linux 4.5




Linux 4.9




Automation

Netflix Vector



Netflix Vector



Netflix Vector


- Open source, on-demand, instance analysis tool
 - <https://github.com/netflix/vector>
- Shows various real-time metrics
- Flame graph support currently in development
 - Automating previous steps
 - Using it internally already
 - Also developing a new d3 front end



Netflix Open Source Software

Advanced Analysis

Linux perf_events Coverage



Advanced Flame Graphs

- Any event can be flame graphed, provided it is issued in *synchronous* Java context
 - Java thread still on-CPU, and event is directly triggered
 - On-CPU Java context is valid
- Synchronous examples:
 - Disk I/O requests issued directly by Java → yes
 - direct reads, sync writes, page faults
 - Disk I/O completion interrupts → no*
 - Disk I/O requests triggered async, e.g., readahead → no*
 - * can be made yes by tracing and associating context

Page Faults

- Show what triggered main memory (resident) to grow:


```
# perf record -e page-faults -p PID -g -- sleep 120
```

- "fault" as (physical) main memory is allocated on-demand, when a virtual page is first populated
- Low overhead tool to solve some types of memory leak

RES column in top(1)

VIRT	RES	COMMAND
3972756	376876	java
344752	231344	ab
0	0	kworker/1:1
1069716	44032	evolution-calen
0	0	ksoftirqd/2

grows
because




Context Switches


- Show why Java blocked and stopped running on-CPU:

```
# perf record -e context-switches -p PID -g -- sleep 5
```

- Identifies locks, I/O, sleeps
 - If code path shouldn't block and looks random, it's an involuntary context switch. I could filter these, but you should have solved them beforehand (CPU load).
- e.g., was used to understand framework differences:



vs




Disk I/O Requests

- Shows who issued disk I/O (sync reads & writes):

```
# perf record -e block:block_rq_insert -a -g -- sleep 60
```

- e.g.: page faults in GC? This JVM has swapped out!:




TCP Events

- TCP transmit, using dynamic tracing:

```
# perf probe tcp_sendmsg
# perf record -e probe:tcp_sendmsg -a -g -- sleep 1; jmaps
# perf script -f comm,pid,tid,cpu,time,event,ip,sym,dso,trace > out.stacks
# perf probe --del tcp_sendmsg
```

- Note: can be high overhead for high packet rates
 - For the current perf trace, dump, post-process cycle
- Can also trace TCP connect & accept (lower overhead)
- TCP receive is async
 - Could trace via socket read




CPU Cache Misses

- In this example, sampling via Last Level Cache loads:

```
# perf record -e LLC-loads -c 10000 -a -g -- sleep 5; jmaps  
# perf script -f comm,pid,tid,cpu,time,event,ip,sym,dso > out.stacks
```



- c is the count (samples once per count)
- Use other CPU counters to sample hits, misses, stalls



CPI Flame Graph


- Cycles Per Instruction!
 - red == instruction heavy
 - blue == cycle heavy (likely mem stall cycles)

zoomed:



Java Package Flame Graph

- Sample on-CPU instruction pointer only (no stack)
 - Don't need -XX:+PreserveFramePointer
- y-axis: package name hierarchy
 - java / util / ArrayList / ::size



Linux 2.6+ (pre-BPF):

```
# perf record -F 199 -a -- sleep 30; ./jmaps
# perf script | ./pkgsplit-perf.sh | ./flamegraph.pl > out.svg
```

no -g (stacks)

Links & References

- Flame Graphs
 - <http://www.brendangregg.com/flamegraphs.html>
 - <http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html>
 - <http://queue.acm.org/detail.cfm?id=2927301>
 - "The Flame Graph" CACM, Vol. 56, No. 6 (June 2016)
 - <http://techblog.netflix.com/2015/07/java-in-flames.html>
 - <http://techblog.netflix.com/2016/04/saving-13-million-computational-minutes.html>
 - <http://techblog.netflix.com/2014/11/nodejs-in-flames.html>
 - <http://www.brendangregg.com/blog/2014-11-09/differential-flame-graphs.html>
- Linux perf_events
 - https://perf.wiki.kernel.org/index.php/Main_Page
 - <http://www.brendangregg.com/perf.html>
 - <http://www.brendangregg.com/blog/2015-02-27/linux-profiling-at-netflix.html>
 - Linux 4.5: <http://www.brendangregg.com/blog/2016-04-30/linux-perf-folded.html>
- Netflix Vector
 - <https://github.com/netflix/vector>
 - <http://techblog.netflix.com/2015/04/introducing-vector-netflixs-on-host.html>
- hprof: <http://www.brendangregg.com/blog/2014-06-09/java-cpu-sampling-using-hprof.html>

Thanks

- Questions?
- <http://techblog.netflix.com>
- <http://slideshare.net/brendangregg>
- <http://www.brendangregg.com>
- bgregg@netflix.com
- [@brendangregg](https://twitter.com/brendangregg)

